

TP 1 - Apprentissage Statistique Appliqué

Nokri Amale, Rahis Erwan, Vuillemot Bertrand

2020 - 2021

Contents

1 **Partie 1** 2

1.1 Cross-Validation with GridSearchCV 2

1.2 Visualizing errors 3

1.3 Changing the loss function 4

1 Partie 1

1.1 Cross-Validation with GridSearchCV

Question: Explain in your report what happens when we run `clf.fit(X_train, Y_train)`. What is the complexity for each of the three following cases?

The line `clf.fit(X_train, Y_train)` here uses the fit method on the object `clf` and taking the train sample. We give the features `X` and the outputs `Y`. The object `clf` is from the class `GridSearchCV` which allows us to find the best hyperparameters among a list we chose. It is taking as parameter an object named `knn` of the class `KNeighborsClassifier()`, a dictionary named `parameters` containing the number of neighbors to be tested in the `knn` algorithm (1 to 5 here) and the `cv` parameter referring to the number of folds to be used in the cross-validation. Basically it will perform a 3-folds cross-validation on a `kNN` model with 1 to 5 neighbors on the train sample and it will allow us to keep the best model. The `kNN` algorithm is parametered with the default metric which is the Euclidean distance : $\sqrt{\sum_{i=1}^n (x_i - y_i)^2}$. The functions are all part of the `sklearn` package.

Question: What is the test accuracy? What would be the accuracy of random guess?

The test accuracy is the measure of how often the points are correctly classified. In our case the accuracy is 0.875. It means that 87.5% of the time, the points are correctly classified on the test sample. If we did a random guess we would randomly choose an output in the range 0 to 9 so the accuracy would converge towards $\frac{1}{10}$ according to the LLN.

Question: What is `LinearSVC()` classifier? Which kernel are we using? What is `C`? (this is a tricky question, try to find the answer online)

`LinearSVC` means Linear Support Vector Classification. We are using a linear kernel. The parameter `C` represents the regularization weights, ie the penalty applied on the loss function. The loss function used here is the Squared Hinge Loss : $l(y) = \max(0, 1 - t \cdot y)$




Figure 1: Example SVM

Question : What is the outcome of `np.logspace(-8, 8, 17, base=2)`? More generally, what is the outcome of `np.logspace(-a, b, k, base=m)`?

The outcome of `np.logspace(-8, 8, 17, base=2)` is a logarithmic space going from 2^{-8} to 2^8 with 17 numbers equally spaced on log scale. The `logspace` function from the `numpy` package will return `k` numbers going from m^{-a} to m^b spaced on a log scale with a log base `m`.

Question : What is the meaning of the warnings? What is the parameter responsible for its appearance?

The warning tells us that the algorithm did not converge, it did not reach the stop criterion. The parameter responsible for its appearance is the `max_iter` parameter. Its value is not sufficient for the algorithm to converge. The data variance is maybe too large for the algorithm to efficiently

perform the SVM.

Question : What did we change with respect to the previous run of LinearSVC()?

Answer: We added a parameter MaxAbsScaler() to scale the absolute data between 0 and 1 and thus reduce the variance.

Question : Explain what happens if we execute

Answer: Those lines will execute a SVM with a maxabsscaler parameter but with no C parameter which is by default 1.0.

Question : what is the difference between StandardScaler() and MaxAbsScaler()? What are other scaling options available in sklearn?

Answer: StandardScaler will normalize the data : $\frac{x-m}{\sigma}$ with m the mean and σ the standard deviation of data. It differs from StandarScaler because absolute values are mapped in the range [0,1].

Question : Using the previous code as an example achieve test accuracy ≥ 0.9 . You can use any method from sklearn package. Give a mathematical description of the selected method. Explain the range of considered hyperparamers.

Answer: choose an algorithm and test

1.2 Visualizing errors

The Logistic Regression function from skLearn package is able to give us the probabilities of each outcome. The figure 2 shows pictures, outcome, real value and probabilities for the 4 first mistakes, ie the 4 first times the predicted value is not equal to the real value.

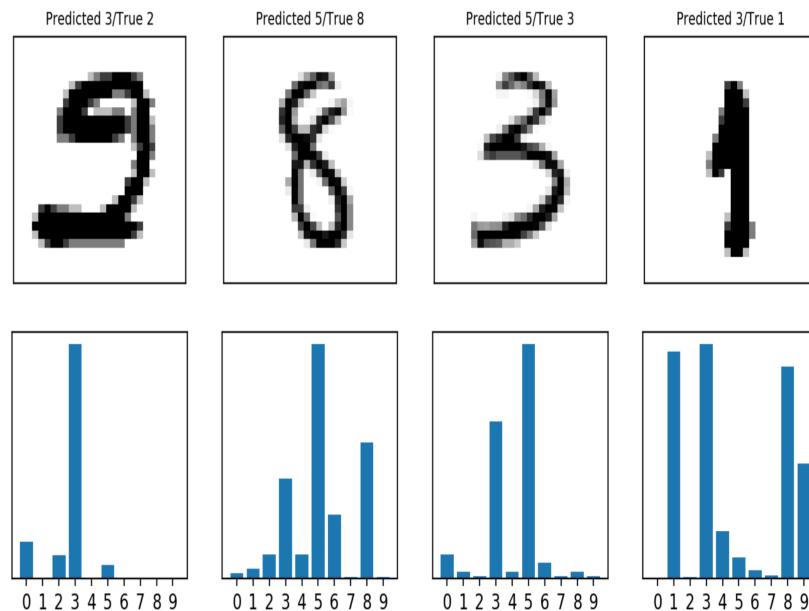


Figure 2: Probabilities of each outcome for the logistic regression

The error in the chunk of code was because the `predict_proba` method returns an array of probabilities within an array. We must then pick the first element of the array (index 0) to obtain the probabilities array (see line 11).

The code is as follows :

```
1 axes = plt.subplots(2, 4)[1] # creates a grid of 10 plots
2
3 # More details about zip() function here https://docs.python.org/3.3/library/
4 # functions.html#zip
5 y_pred = clf4.predict(X_test)
6 j = 0 # Index which iterates over plots
7 for true_label, pred_label, image in list(zip(y_test, y_pred, X_test)):
8     if j == 4: # We only want to look at 4 first mistakes
9         break
10    if true_label != pred_label:
11        # Plotting predicted probabilities
12        axes[1, j].bar(np.arange(10), clf4.predict_proba(image.reshape(1, -1))[0])
13        axes[1, j].set_xticks(np.arange(10))
14        axes[1, j].set_yticks([])
15
16    # Plotting the image
17    axes[0, j].imshow(image.reshape((28, 28)), cmap=plt.cm.gray_r,
18    interpolation='nearest')
19    axes[0, j].set_xticks([])
20    axes[0, j].set_yticks([])
21    axes[0, j].set_title('Predicted {}'.format(pred_label)+' / True {}'.format(
22    true_label), fontsize=8)
23    j += 1
```

1.3 Changing the loss function

Question: What is `balanced_accuracy_score`? Write its mathematical description.

Balanced accuracy score is

References